

COL380 Assignment 1: Sorting with OpenMP Tasks

Read the following sorting algorithm and implement it using OpenMP framework. **You must use OpenMP tasks to implement the algorithm.** Try to parallelize your implementation as much as possible. Implementation is expected to scale with number of CPUs.

Algorithm ParallelSort(A, n, p):

Input:

- A – an array of unsigned integers
- n – number of elements in A
- p – number of buckets

Procedure:

1. Divide A into A_0, A_1, \dots, A_{p-1} , p buckets of size $n/p \pm 1$ each as follows. Each A_i contains contiguous elements of A.
2. From each bucket A_i , select first p elements as pseudo-splitters. Let $R = [r_0, r_1, \dots, r_{p^2-1}]$ be the sorted list of p^2 pseudo-splitters. This sorting may use ParallelSort or SequentialSort.
3. Select $p-1$ equally spaced splitters from R as follows. Let $S = [s_0, s_1, \dots, s_{p-2}]$ be the selected splitters such that $s_j = R[(j+1)*p]$ for j in 0 to p-2.
4. (Using tasks) Split A into p partitions B_0, B_1, \dots, B_{p-1} such that for any element a in partition B_i , $s_{i-1} < a \leq s_i$. Assume $s_{-1} = -\infty$ and $s_p = \infty$.
5. Let n_i denote the number of elements in partition B_i . Sort each partition B_i in a separate task which uses *SequentialSort*(B_i, n_i) if $n_i < \text{Threshold}$, and *ParallelSort*(B_i, n_i, p) otherwise. *SequentialSort* is sequential sorting of your choice implemented in a task.
6. Return concatenation of sorted partitions B_i .

Note that Threshold is minimum number of elements to switch to sequential sorting. Use $\text{Threshold} = 2n/p$, twice the expected size of the partitions. Grading will be in correctness as well as efficiency, scaling up to 24 cores.

Submission Instructions

Submit a single zip file named [Your Entry Number].zip on Moodle with the following:

1. An outline of your OpenMP task-based implementation and design choices made. Explain the degree of your parallelism and scalability (with the help of a graph like below).
2. Graphs of number of CPUs vs. execution time for (at least up to) 24 CPUs and with array of size at least up to 2^{32} .

3. Sources implementing the function signatures provided in the header “psort.h”. Do not include any data files.
4. A makefile that builds a library named “psort” (libsort.a or libsort.so) with the implementation of the functions provided in “psort.h”.

Note

1. You are expected to implement the sorting method ParallelSort().
2. ParallelSort() works in-place. It need not be stable.
3. You are free to define new functions/variables/classes as per your requirement outside specifications of psort.h. Make sure you include those in your submission and in the makefile. Do not change psort.h.
4. You are also given a driver.cpp file which reads the input data from a data file, calls the sorting code, and prints the sorting time. You can use it to run your implementation. You need not include this file in final submission.
5. A sample input data file is provided which has the following format. The first line contains two unsigned integers n and p , respectively followed by n lines containing n unsigned integers.